

Stage-Aware Early Warning of Australian Winter Crop Yield Shortfall Using Daily Weather and Soil Data

Abstract

Early warning of crop yield shortfall is useful only when risk can be updated before the season is effectively complete and when forecast skill is not driven by post-harvest leakage. We study Australian winter crops at state level using 966 region-crop-year yield observations from 1989–2021, daily SILO weather, and Soil and Landscape Grid attributes. Daily weather is converted into May-Jun, May-Jul, May-Aug, May-Sep, and May-Oct feature windows. We separate a no-yield-history weather-soil setting from an operational yield-history setting to distinguish within-season environmental signal from persistent crop-region yield memory. On held-out 2017–2021 seasons, the best no-yield-history May-Oct internal-suite model reaches RMSE 0.821 t/ha and R^2 0.579, while the best operational May-Oct model reaches RMSE 0.660 t/ha and R^2 0.728. Ablation and internal baseline-suite comparisons show that weather and soil provide measurable but secondary monitoring information; persistent crop-region structure and lagged yield history account for much of the strongest operational skill. Risk classification and uncertainty diagnostics support the framework as a state-level monitoring tool for known crop-region histories, while clearly limiting causal, farm-level, and unseen-region claims.

Keywords: crop yield forecasting; early warning; climate risk; Australian agriculture; weather features; soil data; uncertainty

1. Introduction

Australian dryland grain systems face recurring climate risks, including rainfall deficits, heat events, frost, and periods of high evaporative demand. These risks are particularly important for winter crops, where the difference between early vegetative conditions and late-season grain-filling conditions can change the decision value of a forecast. Prior work has shown that Australian wheat yields have been affected by adverse climate trends and by drought-heat exposure [9, 5]. For decision support, however, the operational question is not only whether yield can be estimated after the season, but whether risk can be updated at meaningful within-season stages.

This paper studies state-level early warning for Australian winter crops. The task is to predict yield and low-yield shortfall risk for each region, crop, harvest year, and forecast window. Unlike a post-hoc anomaly attribution or multi-method XAI study, the central question is operational timing: how early can yield-shortfall risk be monitored? A May-Jun estimate supports early monitoring, while a May-Oct estimate approaches a full-season pre-harvest assessment. This framing is aligned with recent early-warning work that emphasizes anticipatory crop-failure monitoring rather than post-season diagnosis alone [1].

The empirical design is intentionally conservative. We remove production and area variables from the main predictors, compute expected-yield and weather-deviation baselines from training years only, and keep the 2017–2021 test period outside all scaling, thresholding, calibration, and residual estimation steps. This avoids turning post-harvest information into apparent forecast skill.

The main contribution is a stage-aware benchmark framework, not a new forecasting algorithm. First, a no-yield-history weather-soil model tests whether daily weather, train-derived weather deviations, and soil background provide measurable evidence of within-season monitoring signal. Second, an operational model adds lagged yield history to quantify how much forecast quality improves when historical production memory is allowed. This distinction matters because a strong operational forecast may partly reflect persistent crop-region productivity rather than purely within-season weather signal.

We make three contributions. First, we formulate Australian winter-crop forecasting as a lead-time-aware shortfall-monitoring benchmark, where predictions are updated from May-Jun to May-Oct rather than evaluated only after the season is nearly complete. Second, we separate a no-yield-history weather-soil setting from an operational yield-history setting, allowing weather-derived early-warning evidence to be distinguished from persistent crop-region productivity memory. Third, we evaluate practical usefulness through temporal holdout testing, internal baseline-suite comparison, feature-regime ablation, stress validation, threshold-tuned risk screening, and uncertainty diagnostics, framing the outputs as state-level decision support rather than causal or farm-level prediction.

We organize the study around four research questions. RQ1 asks how much forecasting skill is available at successive monitoring windows from May-Jun to May-Oct. RQ2 asks how much of this skill comes from stage-aware weather and soil features compared with lagged yield history. RQ3 asks whether low-yield-risk predictions can support screening without becoming automatic decision triggers. RQ4 asks how robust the framework is under rolling-origin, leave-one-crop-out, and leave-one-region-out stress tests.

2. Related Work

2.1. Crop Yield Forecasting and Early Warning

Machine learning has become a common approach for regional crop-yield forecasting, with weather, soil, remote sensing, and management covariates appearing across many systems [24, 18]. At the same time, several studies warn that yield datasets are often small relative to the dimensionality of environmental predictors, making regularized and tree-based tabular models attractive for grain forecasting [15]. Recent work also pushes yield modeling toward multi-modal benchmarks and operational decision-support interfaces [13, 4]. Our setting is smaller and more targeted: state-level Australian winter crops with a fixed May-Oct monitoring calendar.

Early-warning studies differ from generic yield prediction because lead time is part of the objective. A model that works only after most weather information is observed may have limited preparedness value. Preseason and early-season crop-failure forecasts show the importance of evaluating when skill emerges, not only how accurate the final estimate is [1]. We adopt this view by reporting skill separately for each forecast window.

2.2. Stage-Aware Weather and Extreme-Event Features

Weather-yield relationships are stage dependent. Developmental-stage studies show that climatic means and weather extremes can have different effects depending on when they occur during crop growth [22]. Broader climate-risk evidence also emphasizes compound hot-

dry extremes, crop-specific vulnerability, and heterogeneous regional responses [8, 23]. We therefore aggregate daily weather into stage-aware rainfall, heat, cold, evaporative-demand, radiation, and compound-stress features rather than relying only on full-season averages. Recent extreme-weather yield studies further motivate sparse or regularized modeling of drought, heat, flooding, and compound hazards [7, 11].

2.3. Interpretability, Uncertainty, and Decision Support

For agricultural risk monitoring, the output is used to prioritize attention rather than to prove a mechanistic causal pathway. Interpretability and uncertainty are therefore practical requirements, especially when models inform public or business decisions [21, 17]. We use transparent ablation, group permutation importance, and simple lead-time comparisons rather than making the paper a post-hoc XAI study. Explainable boosting and related glassbox tools provide one route for transparent tabular modeling [16], while conformal and conformalized quantile methods motivate distribution-aware intervals with explicit coverage diagnostics [20, 2].

3. Data and Prediction Targets

3.1. Study Panel and Data Sources

The analysis unit is a region-crop-year observation. The yield panel is derived from ABARES Australian Crop Report state-level crop data, which report crop area, production, and yield for major Australian crops at state or territory level [3]. We use the raw ABARES workbook and a derived `yield_panel.csv` after filtering to five winter crops and six state-level regions. Of the 990 possible state-crop-year combinations for 1989–2021, 966 are observed; the missing 24 source entries are summarized in the supplementary material. Table 1 summarizes the modeled panel. Each of the 966 yield observations is expanded over five forecast windows, giving 4,830 window rows. This expansion does not create new yield outcomes: models and metrics are evaluated within forecast windows, and all train/validation/test decisions are made by calendar year so the same region-crop-year outcome cannot be split across evaluation periods.

Table 1: Data summary for the Australian winter crop early-warning panel.

Item	Value
Analysis unit	State/region, crop, harvest year
Yield observations	966
Window-expanded panel rows	4,830
Years	1989–2021
Crops	Barley, Canola, Lupins, Oats, Wheat
Regions	6 Australian states/territories
Forecast windows	May-Jun, May-Jul, May-Aug, May-Sep, May-Oct
Split	train 701, validation 116, test 149

Daily weather is taken from SILO/LongPaddock Australian climate records, which are based on spatial interpolation of observed weather data [10]. The raw project tables retain daily state-level regional series with latitude and longitude metadata; the current benchmark aggregates these regional daily series by forecast cutoff as a state-level exposure proxy and does not claim crop-area-weighted gridded exposure. We use daily rainfall, maximum and minimum temperature, radiation, evaporation, and vapor pressure to build cutoff-specific summaries. Soil covariates are static regional background attributes from ASRIS/Soil and Landscape Grid of Australia [6]. They include available water capacity, bulk density, clay, sand, silt, soil organic carbon, total nitrogen, total phosphorus, pH, cation exchange capacity, depth of soil, and depth of regolith, summarized as topsoil/subsoil or regional aggregates where available. Because these soil summaries are region-level and static, they are treated as background vulnerability covariates, not as farm-level measurements or causal soil interventions.

3.2. Targets

The regression target is yield in tonnes per hectare, denoted $y_{r,c,t}$. We also define an expected-yield baseline from training years and the shortfall

$$s_{r,c,t} = \hat{y}_{r,c,t}^{\text{exp}} - y_{r,c,t}.$$

The expected-yield baseline is fitted using training years only. In the main specification, expected yield is estimated from a crop-region linear trend in calendar year; crop-level linear trends are used when a crop-region series is too short, and the overall training mean is the final tier in this training-only hierarchy. This baseline excludes current-season weather, production, area, lagged yield, and all validation/test outcomes. Validation and test shortfalls are then computed by applying the fixed training-period baseline to later years. The binary low-yield-risk label is one when shortfall exceeds the crop-specific 80th percentile of the training distribution. We use this threshold as a broad watch-list definition rather than as a severe crop-failure definition. This target follows the general idea that yield anomalies or shortfalls require a reference trend before climate impacts can be assessed [14], but here the shortfall is used for early-warning classification rather than post-hoc causal attribution. The classification target is used for secondary risk-alert analyses; the main quantitative results remain regression and ablation metrics.

4. Stage-Aware Features and Models

4.1. Weather Windows

For each forecast window w , daily weather is aggregated only from May through the cutoff month of w . The design intentionally mimics an operational update: May-Jun uses only two months of weather, while May-Oct uses the longest pre-harvest window. No weather feature uses data after the cutoff month. This makes the lead-time comparison interpretable as a sequence of information sets rather than a single full-season model.

The weather feature groups are rainfall accumulation and dry spells, heat and cold exposure, radiation and evaporative demand, and compound hot-dry stress. Rainfall summaries

include totals, rainy days, dry days, short-window maxima, heavy-rain days, and maximum consecutive dry days. Heat/cold features include mean and extreme maximum and minimum temperatures, heat-day thresholds, heat degree days, frost days, and cold days. Energy and dryness features include radiation, evaporation, and vapor pressure summaries. Compound features flag hot-dry days and high-evaporation dry days. These choices are motivated by stage-aware and extreme-weather yield studies [22, 8, 7].

Table 2: Feature groups used for stage-aware early-warning models.

Feature group	Examples	Interpretation
Rainfall/dryness	rain_sum, rain_days, dry_days, dry-spell length	Water supply and dry-spell stress
Heat/cold	tmax_mean, heat_days_30, heat degree days, frost days	Heat exposure and frost/cold risk
Energy/evaporation	radiation_sum, evap_sum, vapor pressure	Growth energy and evaporative demand
Compound stress	hot-dry days, high-evaporation dry days	Combined weather stress
Soil background	AWC, clay/sand, SOC, pH, bulk density, depth	Regional vulnerability conditioning

We additionally construct weather-deviation features by subtracting train-period region-window baselines. This produces rainfall, temperature, dry-spell, heat-day, evapotranspiration, and radiation deviations that are comparable across regions and forecast windows. The weather-deviation baselines are computed using training years only, so they do not leak test-period information.

4.2. No-Yield-History and Operational Feature Sets

We report two primary feature regimes. The no-yield-history weather-soil regime uses stage weather features, train-derived weather-deviation features, soil background features, crop identity, region identity, and year. It excludes lagged yield, rolling yield, production, and area; it is not free of crop, state-level region, or time structure. The operational with yield history regime adds past-yield features from prior years for the same crop-region series. This makes it useful for deployment in known crop-region histories, but it is not interpreted as pure weather evidence.

The model suite combines regularized linear baselines with gradient-boosted tabular models. Ridge and ElasticNet provide stable small-data baselines, while LightGBM and CatBoost provide nonlinear tabular benchmarks [12, 19]. The expanded internal baseline suite also includes Random Forest, SVR-RBF, XGBoost, HistGradientBoosting, GAM, and a daily-weather GRU sequence comparator; implementation details are reported in the supplementary material. We avoid presenting a black-box-only result: every performance table is paired with feature-set ablation, lead-time comparison, or group permutation importance. This is consistent with the broader argument that high-stakes decision support should prefer transparent or at least carefully audited modeling workflows [21].

Classification models use validation-tuned thresholds for low-yield-risk decisions. Instead of fixing a probability threshold at 0.5, we choose thresholds on validation years for F1, recall constraints, or precision constraints, then report test metrics once. Probabilistic interval diagnostics use quantile-style predictions and conformal calibration ideas as supplementary checks [20, 2]. These outputs are not framed as definitive event decisions; they are screening layers for state-level monitoring.

5. Evaluation Protocol

The main split trains on 1989–2012, tunes on 2013–2016, and tests on 2017–2021. The test period is not used for scalars, feature thresholds, residual calibration, conformal residuals, expected-yield baselines, weather-deviation baselines, or low-yield-risk thresholds. This temporal split reflects the intended forecasting direction: models learn from earlier seasons and are evaluated on later seasons.

Regression metrics are MAE, RMSE, and R2. Classification metrics are precision, recall, F1, ROC-AUC, PR-AUC, and Brier score. For probabilistic diagnostics, we summarize interval coverage and width. The conformal intervals target nominal 80% marginal coverage using validation-period calibration residuals only; test coverage is reported without retuning. Deviations from nominal coverage are interpreted as evidence of calibration difficulty under temporal shift and small calibration samples, not as a target corrected using test outcomes. The primary model-selection comparisons are always made within the same forecast window and split.

Robustness checks include rolling-origin folds, leave-one-crop-out validation, and leave-one-region-out validation. Rolling-origin validation tests sensitivity to a single time split. Leave-one-crop-out and leave-one-region-out are intentionally difficult transfer tests. We do not treat them as the main deployment setting, because the intended use case is monitoring crop-region combinations with historical observations. Instead, they identify where the framework is not yet robust enough for unseen-crop or unseen-state generalization.

6. Results

6.1. Lead-Time Skill

Table 3 compares the best no-yield-history weather-soil and operational models at each lead time within the internal baseline suite. For each window and feature regime, the displayed model is selected using validation years only, with RMSE as the selection metric. The operational model is already informative at May-Jun, with RMSE 0.719 t/ha and R2 0.677, and remains strongest at May-Oct with RMSE 0.660 t/ha and R2 0.728. This early operational result may be driven substantially by lagged yield and crop-region memory, so it should not be read as early-season weather-only skill. The no-yield-history suite is weaker but still positive, reaching RMSE 0.821 t/ha and R2 0.579 at May-Oct with an SVR-RBF comparator. Its trajectory is not strictly monotonic: later windows add useful weather information but also correlated features and noise in a small state-level panel. We therefore interpret the curve as a monitoring trajectory rather than assuming that every additional month must improve skill. The gap between regimes motivates the central interpretation of the paper: weather and soil provide measurable but secondary monitoring signal, while historical yield memory is a major driver of operational accuracy. A fixed-model lead-time check is included in the separate supplementary material.

6.2. Ablation: Weather Signal Versus Yield History

Table 4 summarizes the ablation regimes used to isolate information sources. The identity/time baseline is not trivial because crop and state-level region encode persistent productivity differences. Weather-only and weather-plus-soil feature sets add stage-specific en-

Table 3: Validation-selected test performance by forecast window for the no-yield-history weather-soil model and the operational with yield history model.

Window	Feature set	Model	MAE	RMSE	R2
May-Jun	NYH weather-soil	XGBoost	0.651	0.914	0.478
May-Jul	NYH weather-soil	CatBoost	0.659	0.951	0.436
May-Aug	NYH weather-soil	HistGradientBoosting	0.609	0.968	0.416
May-Sep	NYH weather-soil	SVR-RBF	0.609	0.881	0.516
May-Oct	NYH weather-soil	SVR-RBF	0.548	0.821	0.579
May-Jun	Operational	ElasticNet	0.533	0.719	0.677
May-Jul	Operational	Ridge	0.527	0.716	0.680
May-Aug	Operational	ElasticNet	0.506	0.724	0.673
May-Sep	Operational	Ridge	0.469	0.704	0.691
May-Oct	Operational	Ridge	0.428	0.660	0.728




Figure 1: Stage-aware early-warning workflow. Each forecast window uses weather only up to its cutoff month, combines it with static soil and crop-region context, and updates forecasts, risk scores, intervals, and analyst watch lists.

Environmental information, while the full operational model combines this information with past-yield history. The no-yield-history weather-soil models retain positive test skill, especially in the May-Oct window, but they do not consistently outperform identity/time or simple historical baselines. This limits the interpretation of the weather-soil setting: it provides evidence that within-season environmental features contain monitoring signal, not that weather and soil alone dominate persistent crop-region structure. Weather-deviation features are not uniformly beneficial once raw stage summaries and soil are included, which suggests that train-derived deviations should be treated as candidate monitoring features rather than guaranteed improvements. The lag-yield-only feature set is also strong, indicating that persistent crop-region yield history carries important predictive information.




Figure 2: No-yield-history weather-soil model versus operational with yield history model across forecast windows.

Table 4: Feature-regime definitions used for ablation and information-source analysis.

Feature set	Weather	Soil	Yield hist.	Interpretation
Identity/time only	No	No	No	Crop-region persistence benchmark
Weather	Yes	No	No	Stage-weather signal only
Weather+Dev	Yes	No	No	Adds train-derived anomaly features
Weather+Soil	Yes	Yes	No	Adds static regional vulnerability
No-yield-history weather-soil	Yes	Yes	No	Main information-isolation regime
Yield history	No	No	Yes	Persistent productivity memory
Operational with yield history	Yes	Yes	Yes	Main deployment-oriented regime




Figure 3: R2 differences relative to the full operational feature set.

Table 5 adds internal comparator models run under the same temporal split and information constraints. These are not cross-paper SOTA claims, because external crop-yield papers differ in crop, spatial scale, features, and splits. Instead, the table asks whether chronological baselines, classical ML, strong tabular models, interpretable additive models, daily-weather GRU sequence models, and operational yield-history models behave differently on this benchmark. Simple historical baselines remain strong, and the operational models are the main deployment candidates for known crop-region histories. The no-yield-history weather-soil model remains useful as an information-isolation benchmark rather than as a claim of a novel algorithm.

Table 5: Internal baseline-suite comparison under the same leakage-safe split. Each row reports the best validation-selected May-Oct model within the listed group and regime; the comparison is internal to this benchmark and is not a claim of outperforming external crop-yield systems.

Group	Regime	Model	Daily seq.	Yield history	May-Oct RMSE	May-Oct R2
Historical	Past-only yield history	3-year rolling mean	No	Yes	0.779	0.622
Classical ML	No-yield-history weather-soil	SVR-RBF	No	No	0.821	0.579
Strong tabular ML	No-yield-history weather-soil	XGBoost	No	No	0.888	0.508
Interpretable ML	No-yield-history weather-soil	GAM	No	No	1.045	0.318
Sequence	Daily sequence no-yield	DailyWeather-GRU	Yes	No	1.020	0.350
Sequence	Daily sequence + yield history	DailyWeather-GRU	Yes	Yes	0.718	0.678
Strong tabular ML	Operational with yield history	CatBoost	No	Yes	0.778	0.622
Classical ML	Operational with yield history	Ridge	No	Yes	0.660	0.728

6.3. Low-Yield-Risk Classification

Risk classification is treated as a decision-support layer rather than the primary performance claim. With thresholds chosen on validation years only, the best May-Oct classifier is LogisticRegression with threshold 0.27, precision 0.469, recall 0.556, F1 0.508, ROC-AUC 0.797, PR-AUC 0.390, and Brier score 0.135 on the test period. The May-Oct test prevalence is 18.1%, so the PR-AUC is above the random-ranking baseline but still far from a reliable automatic-alert system. The tuned threshold improves recall relative to a naive 0.5 cutoff, which is desirable for early warning where missing a shortfall may be costly. However, the modest precision means alerts should be interpreted as screening signals rather than deterministic event calls. In practice, the classifier would be most useful as a ranked watch list for analyst review, not as an automatic policy trigger.

Table 6: May-Oct low-yield-risk confusion matrix using a validation-selected threshold.

Window	Model	Thr.	TN	FP	FN	TP	Prec.	Recall	F1
May-Oct	LogisticRegression	0.27	105	17	12	15	0.469	0.556	0.508

6.4. Stress Validation

Table 8 summarizes robustness checks using the best model within each validation protocol and feature regime. Values are averaged across the folds and forecast-window rows belonging

Table 7: Held-out May-Oct ranked watch-list example. Watch-list scores are uncalibrated classifier scores used for ranking analyst-review lists and should not be interpreted as fully calibrated event probabilities. Observed indicates the held-out low-yield-risk label.

Rank	Region	Crop	Year	Watch-list score	Alert	Observed	Shortfall
1	VIC	Oats	2018	1.00	1	1	0.465
2	NSW	Lupins	2018	1.00	1	1	0.410
3	VIC	Lupins	2018	1.00	1	0	0.274
4	VIC	Canola	2018	1.00	1	0	0.012
5	NSW	Wheat	2018	0.27	1	1	0.979
6	QLD	Wheat	2019	0.27	1	1	0.957
7	QLD	Wheat	2018	0.27	1	1	0.880
8	QLD	Barley	2019	0.27	1	1	0.865
9	TAS	Oats	2019	0.27	1	1	0.730
10	QLD	Canola	2018	0.27	1	1	0.646

to each protocol, so the held-out-test row is a stress-validation summary rather than a duplicate of the single May-Oct result in Table 3. Rolling-origin performance remains strong, especially for the operational model, suggesting that the headline time-split result is not purely an artifact of one train-test boundary. Leave-one-crop-out is substantially harder for the no-yield-history model, which is expected because crop identity changes the yield scale and sensitivity to weather. Leave-one-region-out produces weak or negative R2. The region-transfer result is the most important limitation: the current framework is better supported for known state-level histories than for unseen-state transfer. This weakness is also consistent with the feature-importance results, where crop-region history and categorical/time structure carry substantial signal.

Table 8: Compact stress validation summary. Rows report mean performance across the folds and forecast windows belonging to each protocol for the best model within a feature regime. The held-out-test row is therefore an all-window stress summary, not the May-Oct-only result in Table 3; leave-one-region-out is a deliberately hard unseen-region test.

Protocol	Feature set	Model	Mean RMSE	Mean R2	Folds
Leave-one-crop-out	Operational with yield history	Ridge	0.731	0.520	5
Leave-one-crop-out	No-yield-history weather-soil	Ridge	1.071	0.067	5
Leave-one-region-out	Operational with yield history	Ridge	0.823	-0.356	6
Leave-one-region-out	No-yield-history weather-soil	LightGBM	0.974	-0.285	6
Rolling-origin	Operational with yield history	Ridge	0.516	0.698	3
Rolling-origin	No-yield-history weather-soil	LightGBM	0.578	0.631	3
Held-out test years, all-window mean	Operational with yield history	Ridge	0.716	0.679	1
Held-out test years, all-window mean	No-yield-history weather-soil	LightGBM	0.963	0.420	1



Figure 4: Stress validation heatmap by validation protocol and feature regime. Values use the same best-model, all-window summary as Table 8, not single-window May-Oct scores.

6.5. Feature Group Importance

Group permutation importance confirms the ablation results. Lag-yield features dominate the operational model, especially in early windows where less current-season weather has accumulated. In the no-yield-history setting, soil background, categorical/time structure, and heat/cold features are among the most important groups. Rainfall and dry-spell features contribute in some windows, but their importance is less stable than the combined history signal. We interpret these patterns as evidence for monitoring and vulnerability conditioning rather than as causal attribution: the model can identify useful statistical signals, but the design is not an intervention study.




Figure 5: Feature group importance measured as mean RMSE increase after group permutation. Operational importance should not be interpreted as within-season weather importance because lag-yield features are allowed in that regime.

7. Discussion

The results suggest a two-tier interpretation. The no-yield-history weather-soil model establishes that stage-aware weather and soil information contain measurable but secondary early-warning signal. The operational model, however, is the more accurate forecasting tool because it also uses historical yield memory. This distinction is practically important: agencies or supply-chain users with stable regional histories may prefer the operational model, while scientific analyses of within-season climate signal should emphasize the no-yield-history model.

Taken together, the results answer the four research questions as follows. RQ1: useful operational skill is available from May-Jun and improves toward May-Oct, while no-yield-history skill is positive but noisier. RQ2: lagged yield history and crop-region persistence explain much of the strongest accuracy, whereas weather-soil features provide a weaker but interpretable monitoring signal. RQ3: risk classification is useful as a ranked screening layer, not as an automatic trigger. RQ4: rolling-origin checks are encouraging, but leave-one-region-out transfer remains weak, limiting deployment to known crop-region histories.

The May-Jun operational result indicates that useful signal exists before the full growing season is observed. May-Oct provides the strongest final pre-harvest performance. Thus, the framework is best seen as a monitoring trajectory rather than a single fixed-date predictor. A decision maker could inspect May-Jun results as an early watch list, update the same crop-region panel as the season progresses, and use May-Sep or May-Oct results for more confident supply-risk assessment.

The strongest practical use cases are state-level monitoring, drought preparedness, procurement planning, and supply-risk assessment. The model does not decide policy or allocate payments; instead, it helps prioritize which crop-region-year combinations deserve attention. This is also why calibration, threshold tuning, and interval coverage are in-

cluded as supplementary diagnostics. Even when a point forecast is accurate on average, risk communication requires knowing how often alerts are missed, how often intervals cover outcomes, and how much uncertainty remains.

A residual-check analysis further clarifies the claim. When a crop-region-year baseline is fitted first, a residual-only weather-soil model does not beat the direct no-yield-history weather-soil model. This negative result supports the conservative narrative: weather and soil signals are useful, but much state-level predictability is tied to persistent crop-region structure and historical yield memory.

Soil attributes should be interpreted as regional conditioning variables. Their incremental accuracy gains are small and not uniform across feature regimes, but they help represent background vulnerability differences that would otherwise be absorbed entirely by crop-region identifiers.

8. Limitations

The analysis is state-level, not farm-level. It does not support field-scale prescription, insurance payout decisions, or causal claims about heat, drought, or soil mechanisms. The soil features are regional background covariates and should not be interpreted as independently manipulated causal factors. Leave-one-region-out validation is weak, so transfer to unseen states remains an open problem. Crop-specific phenological calendars are also left for future work; this study uses a common May-Oct winter-crop proxy for all crops. Finally, the yield panel is small by modern machine-learning standards, so future work should test whether remote sensing, crop calendars, and finer spatial labels can improve transfer without introducing leakage.

9. Conclusion

We presented a stage-aware early-warning benchmark framework for Australian winter crop yield shortfall. The best operational May-Oct model reached RMSE 0.660 t/ha and R2 0.728, while the best no-yield-history weather-soil May-Oct internal-suite model reached RMSE 0.821 t/ha and R2 0.579. The gap between these regimes shows that weather and soil provide measurable but secondary monitoring signal, while historical yield memory is central to the strongest operational forecasts. This supports cautious state-level yield-risk monitoring while clearly bounding farm-level, causal, and unseen-region claims. The framework should be recalibrated before use in unseen regions or states.

References

- [1] Weston Anderson, Shraddhanand Shukla, Jim Verdin, Andrew Hoell, Christina Justice, Brian Barker, Kimberly Slinski, Nathan Lenssen, Jiale Lou, Benjamin I. Cook, and Amy McNally. Preseason maize and wheat yield forecasts for early warning of crop failure. *Nature Communications*, 15(1):1–11, 2024. doi: 10.1038/s41467-024-51555-8.
- [2] Anastasios N. Angelopoulos and Stephen Bates. Conformal prediction: A gentle introduction. *Foundations and Trends in Machine Learning*, 16(4):494–591, 2023. doi: 10.1561/22000000101.

- [3] Australian Bureau of Agricultural and Resource Economics and Sciences. Australian crop report. Australian Government Department of Agriculture, Fisheries and Forestry, 2026. URL <https://www.agriculture.gov.au/abares/research-topics/agricultural-outlook/australian-crop-report>. Accessed 2026-06-24.
- [4] Djavan De Clercq and Adam Mahdi. Feasibility of machine learning-based rice yield prediction in india at the district level using climate reanalysis and remote sensing data. *Agricultural Systems*, 220:104099, 2024. doi: 10.1016/j.agsy.2024.104099.
- [5] Puyu Feng, Bin Wang, Dongliang L. Liu, Cathy Waters, Di Xiao, Li Shi, and Qiang Yu. From climatic risks to impacts: Drought and heat reduce wheat yield in australia. *Agricultural Systems*, 173:303–316, 2019. doi: 10.1016/j.agsy.2019.03.015.
- [6] Mike J. Grundy, Raphael A. Viscarra Rossel, Ross D. Searle, Peter L. Wilson, Changming Chen, and Lisa J. Gregory. Soil and landscape grid of australia. *Soil Research*, 53(8):835–844, 2015. doi: 10.1071/SR15191.
- [7] Jasmin Heilemann, Christian Klassert, Luis Samaniego, Stephan Thober, Andreas Marx, Friedrich Boeing, Bernd Klauer, and Erik Gawel. Projecting impacts of extreme weather events on crop yields using lasso regression. *Weather and Climate Extremes*, 46:100738, 2024. doi: 10.1016/j.wace.2024.100738.
- [8] Matias Heino, Pekka Kinnunen, Weston Anderson, Deepak K. Ray, Michael J. Puma, Olli Varis, Stefan Siebert, and Matti Kummu. Increased probability of hot and dry weather extremes during the growing season threatens global crop yields. *Scientific Reports*, 13:3583, 2023. doi: 10.1038/s41598-023-29378-2.
- [9] Zvi Hochman, David L. Gobbett, and Heidi Horan. Climate trends account for stalled wheat yields in australia since 1990. *Global Change Biology*, 23(5):2071–2081, 2017. doi: 10.1111/gcb.13604.
- [10] S. J. Jeffrey, J. O. Carter, K. B. Moodie, and A. R. Beswick. Using spatial interpolation to construct a comprehensive archive of australian climate data. *Environmental Modelling & Software*, 16(4):309–330, 2001. doi: 10.1016/S1364-8152(01)00008-1.
- [11] Abebe K. Kabtuh and Cheng Qian. Risk of successive hot-pluvial extremes on crop yield loss over global breadbasket regions. *Communications Earth & Environment*, 6: 1040, 2025. doi: 10.1038/s43247-025-02989-5.
- [12] Guolin Ke, Qi Meng, Thomas Finley, Taifeng Wang, Wei Chen, Weidong Ma, Qiwei Ye, and Tie-Yan Liu. Lightgbm: A highly efficient gradient boosting decision tree. In *Advances in Neural Information Processing Systems*, volume 30, pages 3146–3154, 2017. URL <https://papers.nips.cc/paper/6907-lightgbm-a-highly-efficient-gradient-boosting-decision-tree>.
- [13] Fudong Lin, Kaleb Guillot, Summer Crawford, Yihe Zhang, Xu Yuan, and Nian-Feng Tzeng. An open and large-scale dataset for multi-modal climate change-aware crop yield predictions. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 5375–5386, 2024. doi: 10.1145/3637528.3671536.

- [14] Huayue Meng and Long Qian. Performances of different yield-detrending methods in assessing the impacts of agricultural drought and flooding: A case study in the middle-and-lower reach of the yangtze river, china. *Agricultural Water Management*, 296:108812, 2024. doi: 10.1016/j.agwat.2024.108812.
- [15] Michele Meroni, François Waldner, Lucia Seguni, Hervé Kerdiles, and Felix Rembold. Yield forecasting with machine learning and small data: What gains for grains? *Agricultural and Forest Meteorology*, 308–309:108555, 2021. doi: 10.1016/j.agrformet.2021.108555.
- [16] Harsha Nori, Samuel Jenkins, Paul Koch, and Rich Caruana. Interpretml: A unified framework for machine learning interpretability, 2019. URL <https://arxiv.org/abs/1909.09223>.
- [17] Dilli Paudel, Allard de Wit, Hendrik Boogaard, Diego Marcos, Sjoukje Osinga, and Ioannis N. Athanasiadis. Interpretability of deep learning models for crop yield forecasting. *Computers and Electronics in Agriculture*, 206:107663, 2023. doi: 10.1016/j.compag.2023.107663.
- [18] Dilli R. Paudel, Hendrik Boogaard, Allard de Wit, Sander Janssen, Sjoukje Osinga, Christos Pylaniadis, and Ioannis N. Athanasiadis. Machine learning for large-scale crop yield forecasting. *Agricultural Systems*, 187:103016, 2021. doi: 10.1016/j.agsy.2020.103016.
- [19] Liudmila Prokhorenkova, Gleb Gusev, Aleksandr Vorobev, Anna Veronika Dorogush, and Andrey Gulin. Catboost: unbiased boosting with categorical features. In *Advances in Neural Information Processing Systems*, volume 31, pages 6639–6649, 2018. URL <https://papers.nips.cc/paper/7898-catboost-unbiased-boosting-with-categorical-features>.
- [20] Yaniv Romano, Evan Patterson, and Emmanuel Candès. Conformalized quantile regression. In *Advances in Neural Information Processing Systems*, volume 32, 2019. URL <https://arxiv.org/abs/1905.03222>.
- [21] Cynthia Rudin. Why black box machine learning should be avoided for high-stakes decisions, in brief. *Nature Reviews Methods Primers*, 2(1):81, 2022. doi: 10.1038/s43586-022-00172-0.
- [22] Florian Schierhorn, Max Hofmann, Taras Gagalyuk, Igor Ostapchuk, and Daniel Müller. Machine learning reveals complex effects of climatic means and weather extremes on wheat yields during different plant developmental stages. *Climatic Change*, 169(39), 2021. doi: 10.1007/s10584-021-03272-0.
- [23] H. Sjulgard, T. Keller, G. Garland, and T. Colombi. Relationships between weather and yield anomalies vary with crop type and latitude in sweden. *Agricultural Systems*, 211:103757, 2023. doi: 10.1016/j.agsy.2023.103757.
- [24] Thomas van Klompenburg, Ayalew Kassahun, and Cagatay Catal. Crop yield prediction using machine learning: A systematic literature review. *Computers and Electronics in Agriculture*, 177:105709, 2020. doi: 10.1016/j.compag.2020.105709.

