

Stage-Aware Early Warning of Australian Winter Crop Yield Shortfall Using Daily Weather and Soil Data

Abstract

Early warning of crop yield shortfall is useful only when risk can be updated before the season is effectively complete and when forecast skill is not driven by post-harvest leakage. We study Australian winter crops at state level using 966 region-crop-year yield observations from 1989–2021, daily SILO weather, and Soil and Landscape Grid attributes. Daily weather is converted into May-Jun, May-Jul, May-Aug, May-Sep, and May-Oct feature windows. We separate a no-yield-history weather-soil setting from an operational yield-history setting to distinguish within-season environmental signal from persistent crop-region yield memory. On held-out 2017–2021 seasons, the best no-yield-history May-Oct internal-suite model reaches RMSE 0.821 t/ha and R^2 0.579, while the best operational May-Oct model reaches RMSE 0.660 t/ha and R^2 0.728. Ablation and internal baseline-suite comparisons show that weather and soil provide measurable but secondary monitoring information; persistent crop-region structure and lagged yield history account for much of the strongest operational skill. Risk classification and uncertainty diagnostics support the framework as a state-level monitoring tool for known crop-region histories, while clearly limiting causal, farm-level, and unseen-region claims.

Keywords: crop yield forecasting; early warning; climate risk; Australian agriculture; weather features; soil data; uncertainty

1. Introduction

Australian dryland grain systems face recurring climate risks, including rainfall deficits, heat events, frost, and periods of high evaporative demand. These risks are particularly important for winter crops, where the difference between early vegetative conditions and late-season grain-filling conditions can change the decision value of a forecast. Prior work has shown that Australian wheat yields have been affected by adverse climate trends and by drought-heat exposure [9, 5]. For decision support, however, the operational question is not only whether yield can be estimated after the season, but whether risk can be updated at meaningful within-season stages.

This paper studies state-level early warning for Australian winter crops. The task is to predict yield and low-yield shortfall risk for each region, crop, harvest year, and forecast window. Unlike a post-hoc anomaly attribution or multi-method XAI study, the central question is operational timing: how early can yield-shortfall risk be monitored? A May-Jun estimate supports early monitoring, while a May-Oct estimate approaches a full-season pre-harvest assessment. This framing is aligned with early-warning work that emphasizes anticipatory crop-failure monitoring rather than post-season diagnosis alone [1].

The empirical design is intentionally conservative. We remove production and area variables from the main predictors, compute expected-yield and weather-deviation baselines from training years only, and keep the 2017–2021 test period outside all scaling, thresholding, calibration, and residual estimation steps. This avoids turning post-harvest information into apparent forecast skill.

The main contribution is a stage-aware benchmark framework, not a new forecasting algorithm. First, a no-yield-history weather-soil model tests whether daily weather, train-derived weather deviations, and soil background provide measurable evidence of within-season monitoring signal. Second, an operational model adds lagged yield history to quantify how much forecast quality improves when historical production memory is allowed. This distinction matters because a strong operational forecast may partly reflect persistent crop-region productivity rather than purely within-season weather signal.

We make three contributions. First, we formulate Australian winter-crop forecasting as a lead-time-aware shortfall-monitoring benchmark, where predictions are updated from May-Jun to May-Oct rather than evaluated only after the season is nearly complete. Second, we separate a no-yield-history weather-soil setting from an operational yield-history setting, allowing weather-derived early-warning evidence to be distinguished from persistent crop-region productivity memory. Third, we evaluate practical usefulness through temporal holdout testing, internal baseline-suite comparison, feature-regime ablation, stress validation, threshold-tuned risk screening, and uncertainty diagnostics, framing the outputs as state-level decision support rather than causal or farm-level prediction. The machine-learning contribution is therefore the leakage-safe temporal evaluation protocol and benchmark design: staged information sets, explicit feature-regime separation, and stress tests for Australian winter-crop monitoring.

2. Related Work

2.1. Crop Yield Forecasting and Early Warning

Machine learning is now common in regional crop-yield forecasting, where weather, soil, remote sensing, and management covariates appear across many systems [24, 18]. Because yield panels are often small relative to environmental feature sets, regularized and tree-based tabular models remain strong grain-forecasting baselines [15]. Recent multi-modal and decision-support benchmarks broaden this landscape [13, 4]; our setting is narrower: state-level Australian winter crops with a fixed May-Oct monitoring calendar.

Early-warning studies differ from generic yield prediction because lead time is part of the objective. A model that works only after most weather information is observed may have limited preparedness value, so pre-season and early-season crop-failure forecasts evaluate when skill emerges as well as final accuracy [1]. We adopt this view by reporting skill separately for each forecast window.

2.2. Stage-Aware Weather and Extreme-Event Features

Weather-yield relationships are stage dependent. Developmental-stage studies show that climatic means and weather extremes can have different effects depending on when they occur during crop growth [22]. Broader climate-risk evidence emphasizes compound hot-dry extremes, crop-specific vulnerability, heterogeneous regional responses, and sparse or regularized modeling of drought, heat, flooding, and compound hazards [8, 23, 7, 11]. This motivates stage-aware rather than full-season weather summaries.

2.3. Interpretability, Uncertainty, and Decision Support

For agricultural risk monitoring, the output is used to prioritize attention rather than to prove a mechanistic causal pathway. Interpretability and uncertainty are therefore practical requirements, especially when models inform public or business decisions [21, 17]. We use transparent ablation, group permutation importance, and simple lead-time comparisons rather than making the paper a post-hoc XAI study. Explainable boosting and related glassbox tools provide one route for transparent tabular modeling [16], while conformal and conformalized quantile methods motivate distribution-aware intervals with explicit coverage diagnostics [20, 2].

Together, these strands motivate an evaluation-centered benchmark rather than a single-algorithm comparison. The design below therefore emphasizes chronological leakage control, staged information sets, explicit separation between weather-soil evidence and yield-history memory, and diagnostic checks for classification and uncertainty.

3. Methodology

3.1. Study Setting and Data

The analysis unit is a region-crop-year observation. The yield panel is derived from the ABARES Australian Crop Report state crop data workbook `03_AustCropRpt20260303_StateCropData_v1.0.0.xlsx`, accessed on 4 July 2026 [3]. After filtering to five winter crops and six state-level regions, the panel contains 966 observed entries for harvest years 1989–2021 (Figure 1; Table 1). Each yield observation is expanded over five forecast windows, giving 4,830 window rows. This expansion does not create new yield outcomes: models and metrics are evaluated within forecast windows, and all train/validation/test decisions are made by calendar year so the same region-crop-year outcome cannot be split across evaluation periods.

Daily weather is taken from SILO/LongPaddock Australian climate records, which are based on spatial interpolation of observed weather data [10]. The raw project tables retain daily state-level regional series with latitude and longitude metadata; the current benchmark aggregates these regional daily series by forecast cutoff as a state-level exposure proxy and does not claim crop-area-weighted gridded exposure. We use daily rainfall, maximum and minimum temperature, radiation, evaporation, and vapor pressure to build cutoff-specific summaries. Soil covariates are static regional background attributes from ASRIS/Soil and Landscape Grid of Australia [6]. They include available water capacity, bulk density, clay, sand, silt, soil organic carbon, total nitrogen, total phosphorus, pH, cation exchange capacity, depth of soil, and depth of regolith, summarized as topsoil/subsoil or regional aggregates where available. Because these soil summaries are region-level and static, they are treated as background vulnerability covariates, not as farm-level measurements or causal soil interventions.

State-level study regions

966 observed region-crop-year entries, five winter crops, harvest years 1989-2021




Figure 1: Study regions and panel coverage for the Australian winter-crop benchmark. The panel contains 966 observed region-crop-year entries from six state-level regions, five winter crops, and harvest years 1989–2021. QLD and TAS have fewer observed entries because some crop-region-year source records are missing; the benchmark uses state-level regional weather exposure proxies and does not claim crop-area-weighted gridded exposure.

Table 1: Data summary for the Australian winter crop early-warning panel.

Item	Value
Analysis unit	State/region, crop, harvest year
Yield observations	966
Window-expanded panel rows	4,830
Years	1989–2021
Crops	Barley, Canola, Lupins, Oats, Wheat
Regions	6 Australian states/territories
Forecast windows	May-Jun, May-Jul, May-Aug, May-Sep, May-Oct
Split	train 701, validation 116, test 149

3.2. Prediction Targets

The regression target is yield in tonnes per hectare, denoted $y_{r,c,t}$. We also define an expected-yield baseline from training years and the shortfall

$$s_{r,c,t} = \hat{y}_{r,c,t}^{\text{exp}} - y_{r,c,t}.$$

The expected-yield baseline is fitted using training years only. In the main specification, expected yield is estimated from a crop-region linear trend in calendar year; crop-level linear trends are used when a crop-region series is too short, and the overall training mean is the final tier in this training-only hierarchy. This baseline excludes current-season weather, production, area, lagged yield, and all validation/test outcomes. Validation and test shortfalls are then computed by applying the fixed training-period baseline to later years. The binary low-yield-risk label is one when shortfall exceeds the crop-specific 80th percentile of the training distribution. We use this threshold as a broad watch-list definition rather than as a severe crop-failure definition. This target follows the general idea that yield anomalies or shortfalls require a reference trend before climate impacts can be assessed [14], but here the shortfall is used for early-warning classification rather than post-hoc causal attribution. The classification target is used for secondary risk-alert analyses; the main quantitative results remain regression and ablation metrics.

3.3. Feature Construction and Model Regimes

For each forecast window w , daily weather is aggregated only from May through the cutoff month of w . The design intentionally mimics an operational update: May-Jun uses only two months of weather, while May-Oct uses the longest pre-harvest window. No weather feature uses data after the cutoff month. This makes the lead-time comparison interpretable as a sequence of information sets rather than a single full-season model.

The weather feature groups are rainfall accumulation and dry spells, heat and cold exposure, radiation and evaporative demand, and compound hot-dry stress. These choices are motivated by stage-aware and extreme-weather yield studies [22, 8, 7]. We additionally construct weather-deviation features by subtracting train-period region-window baselines. This produces rainfall, temperature, dry-spell, heat-day, evapotranspiration, and radiation deviations that are comparable across regions and forecast windows. The weather-deviation baselines are computed using training years only, so they do not leak test-period information. Feature groups and ablation regimes are listed in the supplementary material.

We report two primary feature regimes. The no-yield-history weather-soil regime uses stage weather features, train-derived weather-deviation features, soil background features, crop identity, region identity, and year. It excludes lagged yield, rolling yield, production, and area; it is not free of crop, state-level region, or time structure. The operational with yield history regime adds past-yield features from prior years for the same crop-region series. This makes it useful for deployment in known crop-region histories, but it is not interpreted as pure weather evidence.

The model suite combines regularized linear baselines with gradient-boosted tabular models. Ridge and ElasticNet provide stable small-data baselines, while LightGBM and CatBoost provide nonlinear tabular benchmarks [12, 19]. The expanded internal baseline suite also includes Random Forest, SVR-RBF, XGBoost, HistGradientBoosting, GAM, and

a daily-weather GRU sequence comparator; implementation details are reported in the supplementary material. We avoid presenting a black-box-only result: every performance table is paired with feature-set ablation, lead-time comparison, or group permutation importance. This is consistent with the broader argument that high-stakes decision support should prefer transparent or at least carefully audited modeling workflows [21].

Classification models use validation-tuned thresholds for low-yield-risk decisions. Instead of fixing a probability threshold at 0.5, we choose thresholds on validation years for F1, recall constraints, or precision constraints, then report test metrics once. Probabilistic interval diagnostics use quantile-style predictions and conformal calibration ideas as supplementary checks [20, 2]. These outputs are not framed as definitive event decisions; they are screening layers for state-level monitoring.

3.4. Experimental Setup

The main split trains on 1989–2012, tunes on 2013–2016, and tests on 2017–2021. The test period is not used for scalars, feature thresholds, residual calibration, conformal residuals, expected-yield baselines, weather-deviation baselines, or low-yield-risk thresholds. This temporal split reflects the intended forecasting direction: models learn from earlier seasons and are evaluated on later seasons.

Regression metrics are MAE, RMSE, and R^2 . Classification metrics are precision, recall, F1, ROC-AUC, PR-AUC, and Brier score. For probabilistic diagnostics, we summarize interval coverage and width. The conformal intervals target nominal 80% marginal coverage using validation-period calibration residuals only; test coverage is reported without retuning. Deviations from nominal coverage are interpreted as evidence of calibration difficulty under temporal shift and small calibration samples, not as a target corrected using test outcomes. The primary model-selection comparisons are always made within the same forecast window and split.

Robustness checks include rolling-origin folds, leave-one-crop-out validation, and leave-one-region-out validation. Rolling-origin validation tests sensitivity to a single time split. Leave-one-crop-out and leave-one-region-out are intentionally difficult transfer tests. We do not treat them as the main deployment setting, because the intended use case is monitoring crop-region combinations with historical observations. Instead, they identify where the framework is not yet robust enough for unseen-crop or unseen-state generalization.

4. Results

4.1. Lead-Time Skill

The first empirical question is whether the monitoring signal strengthens as the season unfolds, and whether that signal is still visible when yield history is withheld. Table 2 gives a consistent answer: forecast accuracy improves as later-season information becomes available, and the operational setting outperforms the no-yield-history setting across all forecast windows. By May-Oct, the best operational model reaches RMSE 0.660 t/ha and R^2 0.728, while the best no-yield-history model reaches RMSE 0.821 t/ha and R^2 0.579.

The May-Jun operational result is already informative, with RMSE 0.719 t/ha and R^2 0.677, but its interpretation is deliberately bounded. It is not early-season weather-only

Table 2: Validation-selected test performance by forecast window for the no-yield-history weather-soil model and the operational with yield history model.

Window	Feature set	Model	MAE	RMSE	R^2
May-Jun	NYH weather-soil	XGBoost	0.651	0.914	0.478
May-Jul	NYH weather-soil	CatBoost	0.659	0.951	0.436
May-Aug	NYH weather-soil	HistGradientBoosting	0.609	0.968	0.416
May-Sep	NYH weather-soil	SVR-RBF	0.609	0.881	0.516
May-Oct	NYH weather-soil	SVR-RBF	0.548	0.821	0.579
May-Jun	Operational	ElasticNet	0.533	0.719	0.677
May-Jul	Operational	Ridge	0.527	0.716	0.680
May-Aug	Operational	ElasticNet	0.506	0.724	0.673
May-Sep	Operational	Ridge	0.469	0.704	0.691
May-Oct	Operational	Ridge	0.428	0.660	0.728

skill; it combines early weather with crop-region memory. The no-yield-history trajectory is weaker and not strictly monotonic, which is plausible in a small state-level panel where later windows add useful weather information alongside correlated features and noise. We therefore read the lead-time curve as a monitoring trajectory rather than as a guarantee that each additional month must improve every model.

4.2. Ablation and Internal Comparisons

The feature-regime comparison explains why the operational curve is strong. Weather and soil provide a measurable monitoring signal, but the larger source of predictive stability is persistent crop-region structure and lagged yield history. This conclusion is consistent across the ablation comparison and the internal baseline suite: no-yield-history weather-soil models retain positive test skill, but they do not consistently dominate identity/time or simple historical baselines.

Table 3 adds internal comparator models under the same temporal split and information constraints. These rows are not cross-paper SOTA claims, because external crop-yield studies differ in crop, spatial scale, features, and evaluation protocol. Their role is more diagnostic: they show whether chronological baselines, classical ML, strong tabular models, interpretable additive models, daily-weather GRU sequence models, and operational yield-history models tell the same story on this benchmark. They do. Simple historical baselines remain strong, operational models are the main deployment candidates for known crop-region histories, and the no-yield-history weather-soil model is best understood as an information-isolation benchmark.

Table 3: Internal baseline-suite comparison under the same leakage-safe split. Each row reports the best validation-selected May-Oct model within the listed group and regime; the comparison is internal to this benchmark and is not a claim of outperforming external crop-yield systems.

Group	Regime	Model	Daily seq.	Yield history	May-Oct RMSE	May-Oct R^2
Historical	Past-only yield history	3-year rolling mean	No	Yes	0.779	0.622
Classical ML	No-yield-history weather-soil	SVR-RBF	No	No	0.821	0.579
Strong tabular ML	No-yield-history weather-soil	XGBoost	No	No	0.888	0.508
Interpretable ML	No-yield-history weather-soil	GAM	No	No	1.045	0.318
Sequence	Daily sequence no-yield	DailyWeather-GRU	Yes	No	1.020	0.350
Sequence	Daily sequence + yield history	DailyWeather-GRU	Yes	Yes	0.718	0.678
Strong tabular ML	Operational with yield history	CatBoost	No	Yes	0.778	0.622
Classical ML	Operational with yield history	Ridge	No	Yes	0.660	0.728

4.3. Low-Yield-Risk Classification

Risk classification is useful only if it is framed as triage rather than automation. With thresholds chosen on validation years only, the best May-Oct classifier is LogisticRegression with threshold 0.27, precision 0.469, recall 0.556, F1 0.508, ROC-AUC 0.797, PR-AUC 0.390, and Brier score 0.135 on the test period. The May-Oct test prevalence is 18.1%, so the PR-AUC is above the random-ranking baseline but not high enough for automatic alerts. The tuned threshold improves recall relative to a naive 0.5 cutoff, which is appropriate when missing a shortfall is costly, but modest precision means the output should be a ranked analyst watch list rather than a deterministic event call. Full threshold sensitivity, confusion-matrix, and ranked watch-list details are reported in the supplementary material.

4.4. Stress Validation

The stress tests clarify where the benchmark is reliable and where it is not. Rolling-origin validation remains comparatively strong, suggesting that the main temporal holdout result is not driven only by one favorable train-test boundary. Performance declines under harder transfer settings, especially when the model is evaluated on unseen regions. This pattern supports the intended use case of monitoring known crop-region histories while warning against extrapolating the framework to unseen state-level regions without recalibration (Table 4).

Leave-one-crop-out is substantially harder for the no-yield-history model, which is expected because crop identity changes the yield scale and sensitivity to weather. Leave-one-region-out produces weak or negative R^2 . The region-transfer result is the most important limitation: the current framework is better supported for known state-level histories than for unseen-state transfer. A heatmap view of the same stress-validation summary is included in the supplementary material.

Table 4: Compact stress validation summary. Rows report mean performance across the folds and forecast windows belonging to each protocol for the best model within a feature regime. The held-out-test row is therefore an all-window stress summary, not the May-Oct-only result in Table 2; leave-one-region-out is a deliberately hard unseen-region test.

Protocol	Feature set	Model	Mean RMSE	Mean R^2	Folds
Leave-one-crop-out	Operational with yield history	Ridge	0.731	0.520	5
Leave-one-crop-out	No-yield-history weather-soil	Ridge	1.071	0.067	5
Leave-one-region-out	Operational with yield history	Ridge	0.823	-0.356	6
Leave-one-region-out	No-yield-history weather-soil	LightGBM	0.974	-0.285	6
Rolling-origin	Operational with yield history	Ridge	0.516	0.698	3
Rolling-origin	No-yield-history weather-soil	LightGBM	0.578	0.631	3
Held-out test years, all-window mean	Operational with yield history	Ridge	0.716	0.679	1
Held-out test years, all-window mean	No-yield-history weather-soil	LightGBM	0.963	0.420	1

4.5. Feature Group Importance

Group permutation importance provides the same interpretation from a different angle. Lag-yield features dominate the operational model, especially in early windows where less current-season weather has accumulated. In the no-yield-history setting, soil background, categorical/time structure, and heat/cold features are among the most important groups, while rainfall and dry-spell features contribute less consistently. These patterns support monitoring and vulnerability conditioning rather than causal attribution: the model identifies useful statistical signals, but the design is not an intervention study. The full feature-importance figure is kept in the supplementary material to avoid duplicating the main ablation evidence.

5. Discussion

The main lesson is not that one algorithm solves Australian crop early warning. It is that the information regime determines what the result means. The no-yield-history weather-soil model establishes that stage-aware weather and soil information contain measurable but secondary early-warning signal. The operational model is the more accurate forecasting tool because it also uses historical yield memory. This distinction is practically important: agencies or supply-chain users with stable regional histories may prefer the operational model, while scientific analyses of within-season climate signal should emphasize the no-yield-history model.

The May-Jun operational result indicates that useful signal exists before the full growing season is observed. May-Oct provides the strongest final pre-harvest performance. Thus, the framework is best seen as a monitoring trajectory rather than a single fixed-date predictor. A decision maker could inspect May-Jun results as an early watch list, update the same crop-region panel as the season progresses, and use May-Sep or May-Oct results for more confident supply-risk assessment.

The strongest practical use cases are state-level monitoring, drought preparedness, procurement planning, and supply-risk assessment. The model does not decide policy or al-

locate payments; instead, it helps prioritize which crop-region-year combinations deserve attention. This is also why calibration, threshold tuning, and interval coverage are included as supplementary diagnostics. Even when a point forecast is accurate on average, risk communication requires knowing how often alerts are missed, how often intervals cover outcomes, and how much uncertainty remains.

A residual-check analysis further clarifies the claim. When a crop-region-year baseline is fitted first, a residual-only weather-soil model does not beat the direct no-yield-history weather-soil model. This negative result supports the conservative narrative: weather and soil signals are useful, but much state-level predictability is tied to persistent crop-region structure and historical yield memory.

Soil attributes should be interpreted as regional conditioning variables. Their incremental accuracy gains are small and not uniform across feature regimes, but they help represent background vulnerability differences that would otherwise be absorbed entirely by crop-region identifiers.

These boundaries define the proper deployment setting. The analysis is state-level, not farm-level. It does not support field-scale prescription, insurance payout decisions, or causal claims about heat, drought, or soil mechanisms. The soil features are regional background covariates and should not be interpreted as independently manipulated causal factors. Leave-one-region-out validation is weak, so transfer to unseen states remains an open problem. The common May-Oct winter-crop calendar is a practical simplification that does not capture crop-specific phenology. The yield panel is also small by modern machine-learning standards. Future work should test whether remote sensing, crop-specific calendars, finer spatial labels, improved transfer learning, and better uncertainty calibration can improve monitoring without introducing leakage.

6. Conclusion

We presented a stage-aware early-warning benchmark framework for Australian winter crop yield shortfall. The best operational May-Oct model reached RMSE 0.660 t/ha and R^2 0.728, while the best no-yield-history weather-soil May-Oct internal-suite model reached RMSE 0.821 t/ha and R^2 0.579. The gap between these regimes shows that weather and soil provide measurable but secondary monitoring signal, while historical yield memory is central to the strongest operational forecasts. This supports cautious state-level yield-risk monitoring while clearly bounding farm-level, causal, and unseen-region claims. The framework should be recalibrated before use in unseen regions or states.

References

- [1] Weston Anderson, Shraddhanand Shukla, Jim Verdin, Andrew Hoell, Christina Justice, Brian Barker, Kimberly Slinski, Nathan Lenssen, Jiale Lou, Benjamin I. Cook, and Amy McNally. Preseason maize and wheat yield forecasts for early warning of crop failure. *Nature Communications*, 15(1):1–11, 2024. doi: 10.1038/s41467-024-51555-8.
- [2] Anastasios N. Angelopoulos and Stephen Bates. Conformal prediction: A gentle introduction. *Foundations and Trends in Machine Learning*, 16(4):494–591, 2023. doi: 10.1561/2200000101.

- [3] Australian Bureau of Agricultural and Resource Economics and Sciences. Australian crop report: State crop data workbook, 2026. State-level crop statistics workbook 03_AustCropRpt20260303_StateCropData_v1.0.0.xlsx, accessed 4 July 2026.
- [4] Djavan De Clercq and Adam Mahdi. Feasibility of machine learning-based rice yield prediction in india at the district level using climate reanalysis and remote sensing data. *Agricultural Systems*, 220:104099, 2024. doi: 10.1016/j.agsy.2024.104099.
- [5] Puyu Feng, Bin Wang, Dongliang L. Liu, Cathy Waters, Di Xiao, Li Shi, and Qiang Yu. From climatic risks to impacts: Drought and heat reduce wheat yield in australia. *Agricultural Systems*, 173:303–316, 2019. doi: 10.1016/j.agsy.2019.03.015.
- [6] Michael J. Grundy, Raphael A. Viscarra Rossel, Ross D. Searle, Peter L. Wilson, Chen Chen, and Luke J. Gregory. Soil and landscape grid of australia. *Soil Research*, 53(8):835–844, 2015.
- [7] Jasmin Heilemann, Christian Klassert, Luis Samaniego, Stephan Thober, Andreas Marx, Friedrich Boeing, Bernd Klauer, and Erik Gawel. Projecting impacts of extreme weather events on crop yields using lasso regression. *Weather and Climate Extremes*, 46:100738, 2024. doi: 10.1016/j.wace.2024.100738.
- [8] Matias Heino, Pekka Kinnunen, Weston Anderson, Deepak K. Ray, Michael J. Puma, Olli Varis, Stefan Siebert, and Matti Kummu. Increased probability of hot and dry weather extremes during the growing season threatens global crop yields. *Scientific Reports*, 13:3583, 2023. doi: 10.1038/s41598-023-29378-2.
- [9] Zvi Hochman, David L. Gobbett, and Heidi Horan. Climate trends account for stalled wheat yields in australia since 1990. *Global Change Biology*, 23(5):2071–2081, 2017. doi: 10.1111/gcb.13604.
- [10] Stephen J. Jeffrey, John O. Carter, Keith B. Moodie, and Alan R. Beswick. Using spatial interpolation to construct a comprehensive archive of australian climate data. *Environmental Modelling and Software*, 16(4):309–330, 2001.
- [11] Abebe K. Kabtiah and Cheng Qian. Risk of successive hot-pluvial extremes on crop yield loss over global breadbasket regions. *Communications Earth & Environment*, 6:1040, 2025. doi: 10.1038/s43247-025-02989-5.
- [12] Guolin Ke, Qi Meng, Thomas Finley, Taifeng Wang, Wei Chen, Weidong Ma, Qiwei Ye, and Tie-Yan Liu. Lightgbm: A highly efficient gradient boosting decision tree. In *Advances in Neural Information Processing Systems*, volume 30, pages 3146–3154, 2017. URL <https://papers.nips.cc/paper/6907-lightgbm-a-highly-efficient-gradient-boosting-decision-tree>.
- [13] Fudong Lin, Kaleb Guillot, Summer Crawford, Yihe Zhang, Xu Yuan, and Nian-Feng Tzeng. An open and large-scale dataset for multi-modal climate change-aware crop yield predictions. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 5375–5386, 2024. doi: 10.1145/3637528.3671536.
- [14] Huayue Meng and Long Qian. Performances of different yield-detrending methods in assessing the impacts of agricultural drought and flooding: A case study in the middle-and-lower reach of the yangtze river, china. *Agricultural Water Management*, 296:108812, 2024. doi: 10.1016/j.agwat.2024.108812.
- [15] Michele Meroni, François Waldner, Lucia Seguí, Hervé Kerdiles, and Felix Rembold. Yield forecasting with machine learning and small data: What gains for grains? *Agricultural and Forest Meteorology*, 308–309:108555, 2021. doi: 10.1016/j.agrformet.2021.108555.
- [16] Harsha Nori, Samuel Jenkins, Paul Koch, and Rich Caruana. Interpretml: A unified framework for machine learning interpretability, 2019. URL <https://arxiv.org/abs/1909.09223>.
- [17] Dilli Paudel, Allard de Wit, Hendrik Boogaard, Diego Marcos, Sjoukje Osinga, and Ioannis N. Athanasiadis. Interpretability of deep learning models for crop yield forecasting. *Computers and Electronics in Agriculture*, 206:107663, 2023. doi: 10.1016/j.compag.2023.107663.
- [18] Dilli R. Paudel, Hendrik Boogaard, Allard de Wit, Sander Janssen, Sjoukje Osinga, Christos Pylianidis, and Ioannis N. Athanasiadis. Machine learning for large-scale crop yield forecasting. *Agricultural Systems*, 187:103016, 2021. doi: 10.1016/j.agsy.2020.103016.

- [19] Liudmila Prokhorenkova, Gleb Gusev, Aleksandr Vorobev, Anna Veronika Dorogush, and Andrey Gulin. Catboost: unbiased boosting with categorical features. In *Advances in Neural Information Processing Systems*, volume 31, pages 6639–6649, 2018. URL <https://papers.nips.cc/paper/7898-catboost-unbiased-boosting-with-categorical-features>.
- [20] Yaniv Romano, Evan Patterson, and Emmanuel Candès. Conformalized quantile regression. In *Advances in Neural Information Processing Systems*, volume 32, 2019. URL <https://arxiv.org/abs/1905.03222>.
- [21] Cynthia Rudin. Why black box machine learning should be avoided for high-stakes decisions, in brief. *Nature Reviews Methods Primers*, 2(1):81, 2022. doi: 10.1038/s43586-022-00172-0.
- [22] Florian Schierhorn, Max Hofmann, Taras Gagalyuk, Igor Ostapchuk, and Daniel Müller. Machine learning reveals complex effects of climatic means and weather extremes on wheat yields during different plant developmental stages. *Climatic Change*, 169(39), 2021. doi: 10.1007/s10584-021-03272-0.
- [23] H. Sjulgard, T. Keller, G. Garland, and T. Colombi. Relationships between weather and yield anomalies vary with crop type and latitude in sweden. *Agricultural Systems*, 211:103757, 2023. doi: 10.1016/j.agry.2023.103757.
- [24] Thomas van Klompenburg, Ayalew Kassahun, and Cagatay Catal. Crop yield prediction using machine learning: A systematic literature review. *Computers and Electronics in Agriculture*, 177:105709, 2020. doi: 10.1016/j.compag.2020.105709.